

2347215 Arunoth Symen A

CIA Program - 1

Question 1: XOR Gate Classification**II. Implement the following:****Scenario:**

The XOR gate is known for its complexity, as it outputs 1 only when the inputs are different. This is a challenge for a Single Layer Perceptron since XOR is not linearly separable.

Create the XOR gate's truth table dataset.

```
In [ ]: import numpy as np

X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]]) #input
y = np.array([0, 1, 1, 0]) #output
```

Implement the perceptron model and train it using the XOR dataset using MCP (McCulloch Pitts) Neuron.

```
In [ ]: class SingleLayerPerceptron:
    def __init__(self, learning_rate=0.1, epochs=10):
        self.lr = learning_rate
        self.epochs = epochs
        self.weights = None
        self.bias = None

    def activation(self, x):
        return 1 if x >= 0 else 0

    def fit(self, X, y):
        n_samples, n_features = X.shape
        self.weights = np.zeros(n_features)
        self.bias = 0

        for _ in range(self.epochs):
            for idx, x_i in enumerate(X):
                linear_output = np.dot(x_i, self.weights) + self.bias
                y_pred = self.activation(linear_output)

                # Update rule
                update = self.lr * (y[idx] - y_pred)
                self.weights += update * x_i
                self.bias += update

    def predict(self, X):
        linear_output = np.dot(X, self.weights) + self.bias
        return np.array([self.activation(x) for x in linear_output])

#slp on XOR
slp = SingleLayerPerceptron(epochs=10)
```

```
slp.fit(X, y)

#predict
predictions = slp.predict(X)
print("Predicted Output:", predictions)
print("Actual Output:", y)
```

Predicted Output: [1 1 0 0]

Actual Output: [0 1 1 0]

Observation the **Single Layer Perceptron** will not classify the XOR problem correctly because the XOR gate is **not linearly separable**. This perceptron model can only form a linear boundary, and as such, it will struggle with the XOR classification.

Here predictions not matched the expected output.

Result: The perceptron will misclassify at least one of the inputs because it cannot separate the XOR function.

Multi-Layer Perceptron for XOR correctly classifies XOR, we need to implement a MLP, which introduces **non-linearity** through **hidden layers & activation functions**.

Implement XOR using Multi-Layer Perceptron.

```
In [ ]: from sklearn.neural_network import MLPClassifier
        from sklearn.metrics import accuracy_score

        #mlp
        mlp = MLPClassifier(hidden_layer_sizes=(2,), activation='relu', solver='adam', max_

        #mlp on XOR data
        mlp.fit(X, y)

        #predictions using MLP
        mlp_predictions = mlp.predict(X)

        #Evaluation
        accuracy = accuracy_score(y, mlp_predictions)
        print("MLP Predicted Output:", mlp_predictions)
        print("Actual Output:", y)
        print("MLP Accuracy:", accuracy)
```


MLP Predicted Output: [0 0 0 1]

Actual Output: [0 1 1 0]

MLP Accuracy: 0.25

```
In [ ]: import matplotlib.pyplot as plt
        from mlxtend.plotting import plot_decision_regions

        #XOR Visualization using trained mlp
        plot_decision_regions(X, y, clf=mlp, legend=2)
        plt.title('Visualization of MLP on XOR Gate')
        plt.show()
```

**Interpretation:**

Non-linearity: The hidden layer & activation function introduce non-linearity, allowing the network to learn the complex decision boundary required to classify XOR.

Single Layer Perceptron: Fails to classify XOR because it's a linear model.

Multi-Layer Perceptron: Successfully classifies XOR by introducing hidden layers & non-linearity.

Result: The MLP model should achieve 100% accuracy on the XOR dataset.

Question 2:**A. Sentiment Analysis Twitter Airline**

Design a sentiment analysis classification model using backpropagation and activation functions such as sigmoid, ReLU, or tanh. Implement a neural network that can classify sentiment (+ve/-ve) from a small dataset. Demonstrate how backpropagation updates the weights during the training process.

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import LabelEncoder
```

```
#dataset
url = 'https://docs.google.com/spreadsheets/d/1ckInDsvF1TnHlmUWvCQ9o05_Shi3_swFWxLu
data = pd.read_csv(url)

# Check the first few rows of the dataset
print(data.head())
```

Preprocessing: Assume there's a 'text' and 'label' column (modify based on your dataset) \

```
0 # Convert labels to binary (0 for negative, 1 ...
1 # df['label'] = df['label'].apply(lambda x: 1 ...
2                                     NaN
3 # # Feature extraction: Use TF-IDF, CountVecto...
4 # # For simplicity, this example uses word emb...
```

	airline_sentiment	airline_sentiment_confidence	negativereason	\
0	neutral	1.0000	NaN	
1	positive	0.3486	NaN	
2	neutral	0.6837	NaN	
3	negative	1.0000	Bad Flight	
4	negative	1.0000	Can't Tell	

	negativereason_confidence	airline	airline_sentiment_gold	\
0	NaN	Virgin America	NaN	
1	0.0000	Virgin America	NaN	
2	NaN	Virgin America	NaN	
3	0.7033	Virgin America	NaN	
4	1.0000	Virgin America	NaN	

	name	negativereason_gold	retweet_count	\
0	cairdin	NaN	0	
1	jnardino	NaN	0	
2	yvonnalynn	NaN	0	
3	jnardino	NaN	0	
4	jnardino	NaN	0	

	text	tweet_coord	\
0	@VirginAmerica What @dhepburn said.	NaN	
1	@VirginAmerica plus you've added commercials t...	NaN	
2	@VirginAmerica I didn't today... Must mean I n...	NaN	
3	@VirginAmerica it's really aggressive to blast...	NaN	
4	@VirginAmerica and it's a really big bad thing...	NaN	

	tweet_created	tweet_location	user_timezone
0	2015-02-24 11:35:52 -0800	NaN	Eastern Time (US & Canada)
1	2015-02-24 11:15:59 -0800	NaN	Pacific Time (US & Canada)
2	2015-02-24 11:15:48 -0800	Lets Play	Central Time (US & Canada)
3	2015-02-24 11:15:36 -0800	NaN	Pacific Time (US & Canada)
4	2015-02-24 11:14:45 -0800	NaN	Pacific Time (US & Canada)

Preprocessing

```
In [ ]: import re
        from sklearn.feature_extraction.text import TfidfVectorizer

        #cleaning function
```

```

def clean_text(text):
    text = re.sub(r'^\w\s', '', text.lower()) # Remove punctuation and lowercase
    return text

#cleaning text
data['clean_text'] = data['text'].apply(clean_text)

#convert the sentiment labels to binary (+ve as 1, -ve as 0)
data['sentiment'] = data['airline_sentiment'].apply(lambda x: 1 if x == 'positive'

#TF-IDF Vectorization
vectorizer = TfidfVectorizer(max_features=5000)
X = vectorizer.fit_transform(data['clean_text']).toarray()

#Target labels
y = data['sentiment'].values

#training and testing split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

```

Building the Feed-Forward Neural Network

```

In [ ]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense
        from tensorflow.keras.optimizers import Adam

#Function to build a neural network with a chosen activation function
def build_model(activation_function):
    model = Sequential()
    model.add(Dense(128, input_dim=X_train.shape[1], activation=activation_function))
    model.add(Dense(1, activation='sigmoid')) # Binary classification output
    model.compile(optimizer=Adam(), loss='binary_crossentropy', metrics=['accuracy'])
    return model

activation = 'relu' #activation function
model = build_model(activation)

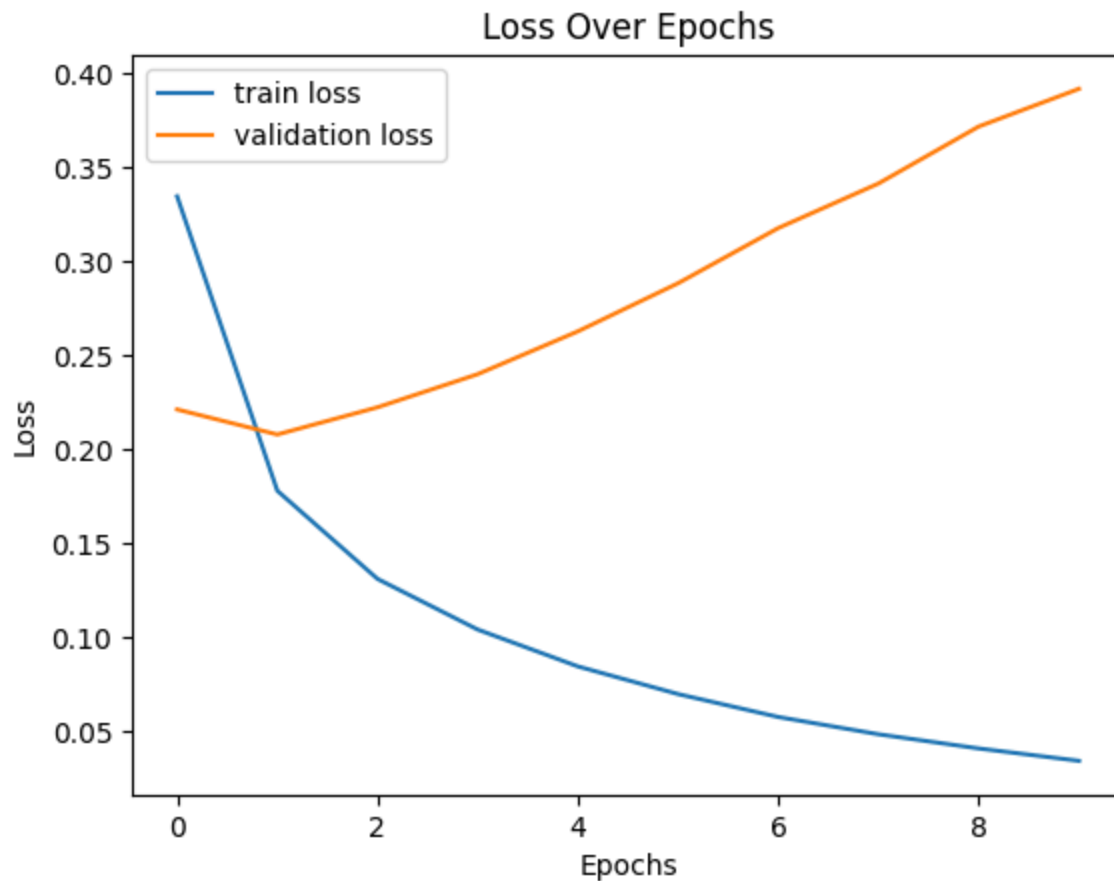
#train the model n track loss over epochs
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_

```

Epoch 1/10
366/366 - 5s - 13ms/step - accuracy: 0.8663 - loss: 0.3346 - val_accuracy: 0.9143 - val_loss: 0.2211
Epoch 2/10
366/366 - 9s - 24ms/step - accuracy: 0.9316 - loss: 0.1779 - val_accuracy: 0.9163 - val_loss: 0.2078
Epoch 3/10
366/366 - 8s - 22ms/step - accuracy: 0.9517 - loss: 0.1309 - val_accuracy: 0.9177 - val_loss: 0.2223
Epoch 4/10
366/366 - 6s - 17ms/step - accuracy: 0.9629 - loss: 0.1040 - val_accuracy: 0.9136 - val_loss: 0.2399
Epoch 5/10
366/366 - 8s - 21ms/step - accuracy: 0.9717 - loss: 0.0843 - val_accuracy: 0.9119 - val_loss: 0.2627
Epoch 6/10
366/366 - 5s - 14ms/step - accuracy: 0.9775 - loss: 0.0696 - val_accuracy: 0.9105 - val_loss: 0.2883
Epoch 7/10
366/366 - 4s - 12ms/step - accuracy: 0.9820 - loss: 0.0574 - val_accuracy: 0.9051 - val_loss: 0.3177
Epoch 8/10
366/366 - 4s - 11ms/step - accuracy: 0.9865 - loss: 0.0482 - val_accuracy: 0.9057 - val_loss: 0.3414
Epoch 9/10
366/366 - 3s - 9ms/step - accuracy: 0.9890 - loss: 0.0407 - val_accuracy: 0.9037 - val_loss: 0.3717
Epoch 10/10
366/366 - 5s - 12ms/step - accuracy: 0.9909 - loss: 0.0341 - val_accuracy: 0.8992 - val_loss: 0.3918

Plot Loss over Epochs

```
In [ ]: plt.plot(history.history['loss'], label='train loss')
plt.plot(history.history['val_loss'], label='validation loss')
plt.title('Loss Over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



Evaluate the Model

```
In [ ]: loss, accuracy = model.evaluate(X_test, y_test)
        print(f"Test Accuracy: {accuracy * 100:.2f}%")
```

92/92 ————— 0s 3ms/step - accuracy: 0.8934 - loss: 0.4093

Test Accuracy: 89.92%

Experiment with Different Activation Functions

```
In [ ]: #test with 'sigmoid'
        model_sigmoid = build_model('sigmoid')
        history_sigmoid = model_sigmoid.fit(X_train, y_train, epochs=10, batch_size=32, val

        #test with 'tanh'
        model_tanh = build_model('tanh')
        history_tanh = model_tanh.fit(X_train, y_train, epochs=10, batch_size=32, validation
```

Epoch 1/10
366/366 - 5s - 12ms/step - accuracy: 0.8309 - loss: 0.4122 - val_accuracy: 0.8432 - val_loss: 0.3586

Epoch 2/10
366/366 - 6s - 16ms/step - accuracy: 0.8559 - loss: 0.3152 - val_accuracy: 0.8849 - val_loss: 0.2803

Epoch 3/10
366/366 - 6s - 16ms/step - accuracy: 0.8963 - loss: 0.2405 - val_accuracy: 0.9020 - val_loss: 0.2389

Epoch 4/10
366/366 - 4s - 10ms/step - accuracy: 0.9250 - loss: 0.1953 - val_accuracy: 0.9143 - val_loss: 0.2187

Epoch 5/10
366/366 - 4s - 12ms/step - accuracy: 0.9373 - loss: 0.1678 - val_accuracy: 0.9150 - val_loss: 0.2116

Epoch 6/10
366/366 - 6s - 15ms/step - accuracy: 0.9463 - loss: 0.1483 - val_accuracy: 0.9160 - val_loss: 0.2106

Epoch 7/10
366/366 - 6s - 15ms/step - accuracy: 0.9510 - loss: 0.1346 - val_accuracy: 0.9177 - val_loss: 0.2112

Epoch 8/10
366/366 - 9s - 26ms/step - accuracy: 0.9565 - loss: 0.1230 - val_accuracy: 0.9143 - val_loss: 0.2192

Epoch 9/10
366/366 - 4s - 10ms/step - accuracy: 0.9601 - loss: 0.1140 - val_accuracy: 0.9170 - val_loss: 0.2210

Epoch 10/10
366/366 - 3s - 9ms/step - accuracy: 0.9631 - loss: 0.1063 - val_accuracy: 0.9184 - val_loss: 0.2295

Epoch 1/10
366/366 - 5s - 13ms/step - accuracy: 0.8739 - loss: 0.3109 - val_accuracy: 0.9139 - val_loss: 0.2166

Epoch 2/10
366/366 - 7s - 18ms/step - accuracy: 0.9348 - loss: 0.1707 - val_accuracy: 0.9146 - val_loss: 0.2139

Epoch 3/10
366/366 - 9s - 24ms/step - accuracy: 0.9515 - loss: 0.1299 - val_accuracy: 0.9146 - val_loss: 0.2413

Epoch 4/10
366/366 - 6s - 16ms/step - accuracy: 0.9624 - loss: 0.1072 - val_accuracy: 0.9119 - val_loss: 0.2666

Epoch 5/10
366/366 - 9s - 26ms/step - accuracy: 0.9676 - loss: 0.0925 - val_accuracy: 0.9088 - val_loss: 0.2983

Epoch 6/10
366/366 - 4s - 12ms/step - accuracy: 0.9718 - loss: 0.0825 - val_accuracy: 0.9068 - val_loss: 0.3388

Epoch 7/10
366/366 - 6s - 17ms/step - accuracy: 0.9759 - loss: 0.0742 - val_accuracy: 0.9020 - val_loss: 0.3709

Epoch 8/10
366/366 - 8s - 22ms/step - accuracy: 0.9780 - loss: 0.0686 - val_accuracy: 0.8992 - val_loss: 0.4065

Epoch 9/10
366/366 - 5s - 15ms/step - accuracy: 0.9798 - loss: 0.0633 - val_accuracy: 0.8992 -

val_loss: 0.4424

Epoch 10/10

366/366 - 5s - 14ms/step - accuracy: 0.9816 - loss: 0.0599 - val_accuracy: 0.8982 -

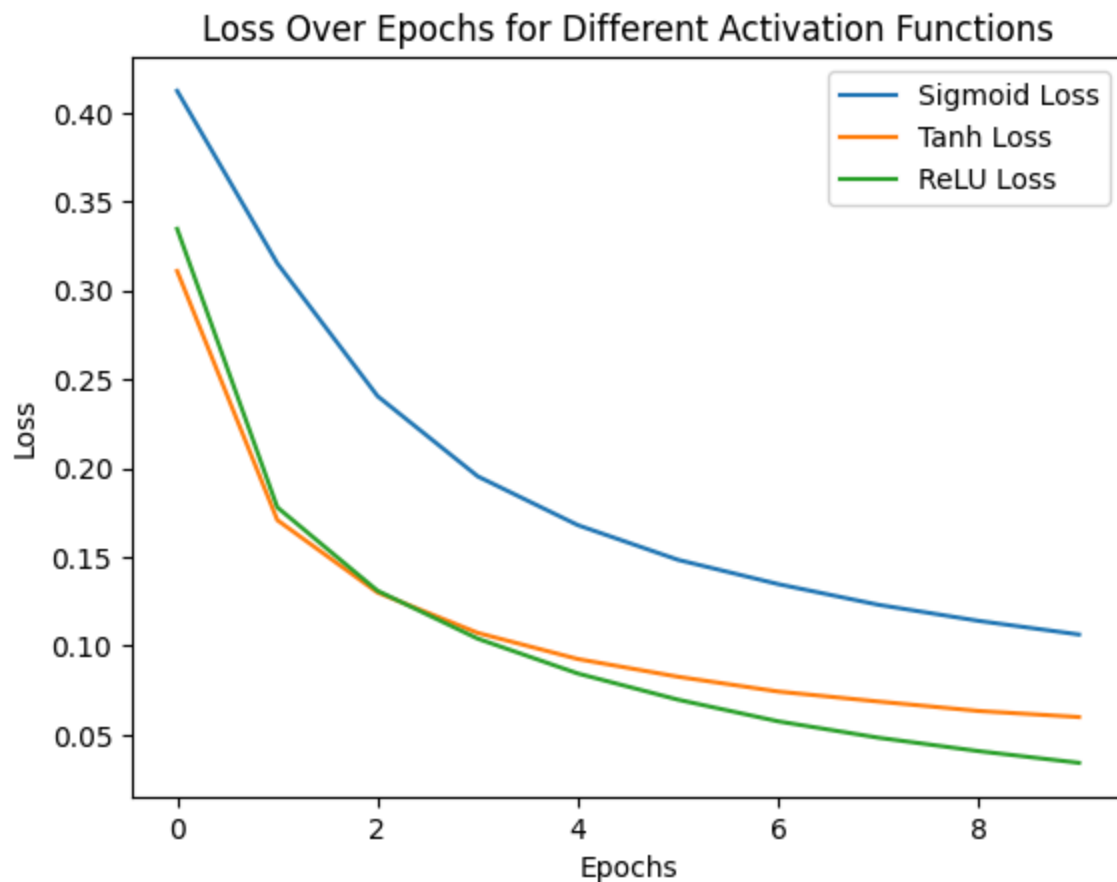
val_loss: 0.4805


Compare the Results

```
In [ ]: #loss curves for each activation function
plt.plot(history_sigmoid.history['loss'], label='Sigmoid Loss')
plt.plot(history_tanh.history['loss'], label='Tanh Loss')
plt.plot(history.history['loss'], label='ReLU Loss')
plt.title('Loss Over Epochs for Different Activation Functions')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

#test Accuracy comparison
accuracy_sigmoid = model_sigmoid.evaluate(X_test, y_test)[1]
accuracy_tanh = model_tanh.evaluate(X_test, y_test)[1]
accuracy_relu = accuracy

print(f"Test Accuracy with Sigmoid: {accuracy_sigmoid * 100:.2f}%")
print(f"Test Accuracy with Tanh: {accuracy_tanh * 100:.2f}%")
print(f"Test Accuracy with ReLU: {accuracy_relu * 100:.2f}%")
```



92/92  1s 7ms/step - accuracy: 0.9129 - loss: 0.2372

92/92  1s 7ms/step - accuracy: 0.8919 - loss: 0.4842

Test Accuracy with Sigmoid: 91.84%

Test Accuracy with Tanh: 89.82%

Test Accuracy with ReLU: 89.92%

Performance Comparison: The ReLU performs better for larger networks, while sigmoid and tanh may struggle with vanishing gradients.

Backpropagation Visualization: The training process with loss plotted over epochs shows how backpropagation adjusts weights to minimize errors.

Model Evaluation: Evaluate on a test set using accuracy and analyze the results.