

2347215 Arunoth Symen A

Lab Program - 3

1. Data Preprocessing

```
In [ ]: from tensorflow.keras.datasets import cifar10
        from tensorflow.keras.utils import to_categorical

        # Load the dataset
        (x_train, y_train), (x_test, y_test) = cifar10.load_data()

        # Normalize the pixel values to [0, 1]
        x_train, x_test = x_train / 255.0, x_test / 255.0

        # Convert class labels to one-hot encoding
        y_train = to_categorical(y_train, 10)
        y_test = to_categorical(y_test, 10)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 ————— 4s 0us/step

Data Augmentation: To improve generalization, apply data augmentation techniques.

```
In [ ]: from tensorflow.keras.preprocessing.image import ImageDataGenerator

        datagen = ImageDataGenerator(
            horizontal_flip=True,
            rotation_range=10,
            width_shift_range=0.1,
            height_shift_range=0.1
        )
        datagen.fit(x_train)
```

2. Network Architecture Design

Architecture Justification:

Input Layer: The CIFAR-10 images have a shape of 32x32x3 (RGB color).

Hidden Layers: • Convolutional layers are used to capture spatial hierarchies and local patterns from the images.

- MaxPooling layers are used for down-sampling, reducing spatial dimensions.
- Dense layers at the end for classification.

Output Layer: Use a softmax activation function to predict the 10 classes.

```
In [ ]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

        model = Sequential()
```

```
# Input Layer with Conv and MaxPooling
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())

# Fully connected layers
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))

model.summary()
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	
conv2d (Conv2D)	(None, 30, 30, 32)	
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	
conv2d_1 (Conv2D)	(None, 13, 13, 64)	
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	
flatten (Flatten)	(None, 2304)	
dense (Dense)	(None, 128)	
dropout (Dropout)	(None, 128)	
dense_1 (Dense)	(None, 10)	



Total params: 315,722 (1.20 MB)

Trainable params: 315,722 (1.20 MB)

Non-trainable params: 0 (0.00 B)

3. Activation Functions

- ReLU (Rectified Linear Unit): Helps to solve the vanishing gradient problem and allows faster training. It's popular in convolutional layers.
- Softmax: Used in the final layer to output class probabilities for multi-class classification.

4. Loss Function and Optimizer

For this classification problem:

- Use categorical crossentropy since it is suited for multi-class classification problems.
- Compare mean squared error and categorical hinge with categorical crossentropy.

```
In [ ]: from tensorflow.keras.losses import categorical_crossentropy
        from tensorflow.keras.optimizers import Adam


        # Compile model with categorical crossentropy and Adam optimizer
        model.compile(optimizer=Adam(learning_rate=0.001),
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```


5. Training the Model


```
In [ ]: history = model.fit(datagen.flow(x_train, y_train, batch_size=64),
                           epochs=50,
                           validation_data=(x_test, y_test))
```


Epoch 1/50


```
/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
  self._warn_if_super_not_called()
```


782/782  **107s** 129ms/step - accuracy: 0.2891 - loss: 1.9154 - val_accuracy: 0.5128 - val_loss: 1.3687
Epoch 2/50


782/782  **129s** 113ms/step - accuracy: 0.4462 - loss: 1.5237 - val_accuracy: 0.5690 - val_loss: 1.2181
Epoch 3/50


782/782  **142s** 114ms/step - accuracy: 0.5029 - loss: 1.3841 - val_accuracy: 0.5837 - val_loss: 1.1997
Epoch 4/50


782/782  **144s** 116ms/step - accuracy: 0.5285 - loss: 1.3178 - val_accuracy: 0.6245 - val_loss: 1.0743
Epoch 5/50


782/782  **141s** 116ms/step - accuracy: 0.5587 - loss: 1.2513 - val_accuracy: 0.6237 - val_loss: 1.0612
Epoch 6/50


782/782  **90s** 115ms/step - accuracy: 0.5700 - loss: 1.2211 - val_accuracy: 0.6498 - val_loss: 1.0108
Epoch 7/50


782/782  **141s** 114ms/step - accuracy: 0.5787 - loss: 1.1920 - val_accuracy: 0.6394 - val_loss: 1.0251
Epoch 8/50


782/782  **143s** 116ms/step - accuracy: 0.5914 - loss: 1.1588 - val_accuracy: 0.6683 - val_loss: 0.9551
Epoch 9/50


782/782  **89s** 114ms/step - accuracy: 0.6007 - loss: 1.1381 - val_accuracy: 0.6697 - val_loss: 0.9527
Epoch 10/50


782/782  **91s** 116ms/step - accuracy: 0.6071 - loss: 1.1152 - val_accuracy: 0.6678 - val_loss: 0.9596
Epoch 11/50


782/782  **89s** 113ms/step - accuracy: 0.6200 - loss: 1.0963 - val_accuracy: 0.6938 - val_loss: 0.8904
Epoch 12/50


782/782  **90s** 115ms/step - accuracy: 0.6246 - loss: 1.0712 - val_accuracy: 0.6746 - val_loss: 0.9408
Epoch 13/50


782/782  **89s** 113ms/step - accuracy: 0.6269 - loss: 1.0660 - val_accuracy: 0.7007 - val_loss: 0.8703
Epoch 14/50


782/782  **89s** 114ms/step - accuracy: 0.6351 - loss: 1.0430 - val_accuracy: 0.6892 - val_loss: 0.9041
Epoch 15/50

782/782  **90s** 115ms/step - accuracy: 0.6349 - loss: 1.0311 - val_accuracy: 0.6968 - val_loss: 0.8652
Epoch 16/50

782/782  **142s** 115ms/step - accuracy: 0.6445 - loss: 1.0155 - val_accuracy: 0.6886 - val_loss: 0.8933
Epoch 17/50

782/782  **88s** 113ms/step - accuracy: 0.6512 - loss: 1.0078 - val_accuracy: 0.6971 - val_loss: 0.8716
Epoch 18/50

782/782  **90s** 115ms/step - accuracy: 0.6507 - loss: 0.9973 - val_accuracy: 0.7067 - val_loss: 0.8509
Epoch 19/50

782/782  **91s** 116ms/step - accuracy: 0.6516 - loss: 1.0012 - val_accuracy: 0.7019 - val_loss: 0.8609

Epoch 20/50
782/782 ————— 142s 116ms/step - accuracy: 0.6580 - loss: 0.9848 - val_accuracy: 0.6988 - val_loss: 0.8683

Epoch 21/50
782/782 ————— 142s 116ms/step - accuracy: 0.6610 - loss: 0.9816 - val_accuracy: 0.7112 - val_loss: 0.8487

Epoch 22/50
782/782 ————— 141s 115ms/step - accuracy: 0.6614 - loss: 0.9735 - val_accuracy: 0.7166 - val_loss: 0.8232

Epoch 23/50
782/782 ————— 142s 115ms/step - accuracy: 0.6681 - loss: 0.9534 - val_accuracy: 0.7098 - val_loss: 0.8229

Epoch 24/50
782/782 ————— 89s 114ms/step - accuracy: 0.6672 - loss: 0.9468 - val_accuracy: 0.7141 - val_loss: 0.8212

Epoch 25/50
782/782 ————— 144s 116ms/step - accuracy: 0.6739 - loss: 0.9450 - val_accuracy: 0.7376 - val_loss: 0.7612

Epoch 26/50
782/782 ————— 91s 116ms/step - accuracy: 0.6668 - loss: 0.9566 - val_accuracy: 0.7281 - val_loss: 0.8048

Epoch 27/50
782/782 ————— 140s 113ms/step - accuracy: 0.6740 - loss: 0.9375 - val_accuracy: 0.7173 - val_loss: 0.8179

Epoch 28/50
782/782 ————— 90s 114ms/step - accuracy: 0.6767 - loss: 0.9394 - val_accuracy: 0.7297 - val_loss: 0.7824

Epoch 29/50
782/782 ————— 142s 114ms/step - accuracy: 0.6743 - loss: 0.9294 - val_accuracy: 0.7275 - val_loss: 0.7874

Epoch 30/50
782/782 ————— 91s 116ms/step - accuracy: 0.6806 - loss: 0.9176 - val_accuracy: 0.7171 - val_loss: 0.8286

Epoch 31/50
782/782 ————— 89s 114ms/step - accuracy: 0.6818 - loss: 0.9181 - val_accuracy: 0.7223 - val_loss: 0.7997

Epoch 32/50
782/782 ————— 91s 116ms/step - accuracy: 0.6854 - loss: 0.9064 - val_accuracy: 0.7381 - val_loss: 0.7556

Epoch 33/50
782/782 ————— 141s 115ms/step - accuracy: 0.6869 - loss: 0.9025 - val_accuracy: 0.7375 - val_loss: 0.7722

Epoch 34/50
782/782 ————— 144s 117ms/step - accuracy: 0.6894 - loss: 0.8964 - val_accuracy: 0.7367 - val_loss: 0.7483

Epoch 35/50
782/782 ————— 141s 116ms/step - accuracy: 0.6931 - loss: 0.8825 - val_accuracy: 0.7152 - val_loss: 0.8165

Epoch 36/50
782/782 ————— 89s 113ms/step - accuracy: 0.6948 - loss: 0.8957 - val_accuracy: 0.7264 - val_loss: 0.7956

Epoch 37/50
782/782 ————— 145s 117ms/step - accuracy: 0.6955 - loss: 0.8871 - val_accuracy: 0.7392 - val_loss: 0.7570

Epoch 38/50
782/782 ————— 90s 115ms/step - accuracy: 0.6970 - loss: 0.8824 - val_

```

accuracy: 0.7477 - val_loss: 0.7344
Epoch 39/50
782/782 ————— 91s 116ms/step - accuracy: 0.6954 - loss: 0.8767 - val_
accuracy: 0.7443 - val_loss: 0.7535
Epoch 40/50
782/782 ————— 92s 117ms/step - accuracy: 0.6940 - loss: 0.8828 - val_
accuracy: 0.7490 - val_loss: 0.7302
Epoch 41/50
782/782 ————— 141s 116ms/step - accuracy: 0.6985 - loss: 0.8776 - val_
accuracy: 0.7342 - val_loss: 0.7706
Epoch 42/50
782/782 ————— 89s 113ms/step - accuracy: 0.6923 - loss: 0.8788 - val_
accuracy: 0.7443 - val_loss: 0.7591
Epoch 43/50
782/782 ————— 91s 116ms/step - accuracy: 0.6932 - loss: 0.8743 - val_
accuracy: 0.7467 - val_loss: 0.7177
Epoch 44/50
782/782 ————— 91s 117ms/step - accuracy: 0.6972 - loss: 0.8711 - val_
accuracy: 0.7476 - val_loss: 0.7366
Epoch 45/50
782/782 ————— 90s 115ms/step - accuracy: 0.6966 - loss: 0.8671 - val_
accuracy: 0.7374 - val_loss: 0.7619
Epoch 46/50
782/782 ————— 142s 114ms/step - accuracy: 0.7016 - loss: 0.8559 - val_
accuracy: 0.7523 - val_loss: 0.7264
Epoch 47/50
782/782 ————— 142s 114ms/step - accuracy: 0.7007 - loss: 0.8574 - val_
accuracy: 0.7520 - val_loss: 0.7237
Epoch 48/50
782/782 ————— 92s 117ms/step - accuracy: 0.7042 - loss: 0.8579 - val_
accuracy: 0.7397 - val_loss: 0.7525
Epoch 49/50
782/782 ————— 141s 116ms/step - accuracy: 0.7047 - loss: 0.8431 - val_
accuracy: 0.7317 - val_loss: 0.7810
Epoch 50/50
782/782 ————— 144s 118ms/step - accuracy: 0.7080 - loss: 0.8409 - val_
accuracy: 0.7594 - val_loss: 0.7135

```

6. Model Evaluation

```
In [ ]: from sklearn.metrics import classification_report, confusion_matrix
```

```

# Evaluate model
y_pred = model.predict(x_test)
y_pred_classes = y_pred.argmax(axis=1)
y_true = y_test.argmax(axis=1)

# Classification report
print(classification_report(y_true, y_pred_classes))

# Confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred_classes)
print(conf_matrix)

```

313/313	6s 18ms/step			
	precision	recall	f1-score	support
0	0.76	0.84	0.80	1000
1	0.82	0.92	0.87	1000
2	0.79	0.57	0.66	1000
3	0.65	0.47	0.55	1000
4	0.75	0.70	0.72	1000
5	0.68	0.64	0.66	1000
6	0.79	0.83	0.81	1000
7	0.69	0.89	0.78	1000
8	0.85	0.85	0.85	1000
9	0.79	0.88	0.83	1000
accuracy				0.76 10000
macro avg				0.76 0.76 0.75 10000
weighted avg				0.76 0.76 0.75 10000

```
[[841 27 13 6 5 1 6 14 40 47]
 [ 10 919 1 2 0 1 5 1 12 49]
 [ 86 9 566 28 88 58 65 66 14 20]
 [ 23 15 48 474 60 164 71 69 34 42]
 [ 26 4 31 25 701 22 42 130 12 7]
 [ 14 6 21 139 30 644 28 88 13 17]
 [ 14 12 28 31 33 19 835 12 7 9]
 [ 16 5 11 12 20 30 3 885 2 16]
 [ 51 52 0 7 0 1 1 5 850 33]
 [ 20 66 1 2 1 1 6 6 18 879]]
```

7. Optimization Strategies

- Early Stopping: Stop training when validation loss stops decreasing to avoid overfitting.
- Learning Rate Scheduling: Gradually reduce the learning rate for smoother convergence.
- Weight Initialization: Proper initialization (e.g., Xavier or He initialization) avoids vanishing/exploding gradients and ensures efficient training.

Weight Initialization Importance Good weight initialization speeds up convergence by ensuring that activations are properly distributed across layers. Without proper initialization, the network can get stuck in a local minimum.

8. Report


```
In [ ]: import matplotlib.pyplot as plt


# Plot the training loss and validation loss
def plot_loss(history):
    plt.figure(figsize=(10, 6))
    plt.plot(history.history['loss'], label='Training Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.title('Loss over Epochs')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()
    plt.grid(True)
```

```
plt.show()

# Plot the training accuracy and validation accuracy
def plot_accuracy(history):
    plt.figure(figsize=(10, 6))
    plt.plot(history.history['accuracy'], label='Training Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title('Accuracy over Epochs')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.grid(True)
    plt.show()

# Visualize
plot_loss(history)
plot_accuracy(history)
```





Here are some of the key elements you would include in a report:

- **Architecture:** Describe and justify the layers, number of filters, and activation functions.
- **Training Results:** Plot loss and accuracy over epochs for training and validation.
- **Hyperparameters:** List the values for learning rate, batch size, and number of epochs.
- **Challenges and Solutions:** Mention challenges like overfitting, slow convergence, or vanishing gradients, and how you addressed them (e.g., by using regularization, adjusting learning rate, or fine-tuning the architecture).